

CSCE 313 — Midterm 2 Cheat Sheet (Comprehensive)

1 Process API & exec Pattern

- Sequence:** `fork() → exec*() → wait()|waitpid() → exit()`
- `execvp()` replaces the process image — it **never returns** if successful.
- Use `waitpid(pid, &status, 0)` to block until child finishes.

Example:

```
pid_t pid = fork();
if (pid == 0) {
    char *args[] = {"ls", "-l", NULL};
    execvp(args[0], args);
    perror("execvp");
    _exit(127);
}
waitpid(pid, NULL, 0);
```

2 Unix File Descriptors (FDs)

- Standard FDs:** `0 = stdin, 1 = stdout, 2 = stderr`
- `open()` returns the **lowest unused fd**.
- File offset and flags live in the **open-file table**.
- `dup2(fd, newfd)` duplicates FDs → used for redirection.
- `close(fd)` releases file descriptors to prevent leaks.

Redirect stdout → file

```
int fd = open("out.txt", O_CREAT|O_WRONLY|O_TRUNC, 0644);
dup2(fd, 1);
close(fd);
execvp(argv[0], argv);
```

3 Common open() Flags

Flag	Meaning	Typical Use
O_RDONLY	Open for read only	reading existing files
O_WRONLY	Open for write only	write logs, append output
O_RDWR	Read and write	random access files

1 / 7

Example: `open("log.txt", O_CREAT|O_WRONLY|O_APPEND, 0644);`

4 Descriptor, File, and Inode Tables

Table	Scope	Contains
FD table	Per-process	integers (0,1,2,...) → open-file entries
Open-file table	Kernel-wide	file offset, flags, refcount
Inode/Vnode	System-wide	metadata (mode, size, owner), block pointers

Multiple processes can share an open-file entry (e.g., after `fork()`), but each has its own FD table.

5 Pipes & Pipelines

- `pipe(fd)` → creates unidirectional data channel.
- `p[0]` = read end, `p[1]` = write end.
- Close unused ends in both parent & child.
- Used for connecting processes via stdin/stdout.

Example (`ls | wc -l`):

```
int p[2]; pipe(p);
if (fork()==0){ dup2(p[1],1); close(p[0]); execvp("ls","ls",NULL); }
if (fork()==0){ dup2(p[0],0); close(p[1]); execvp("wc","wc","-l",NULL); }
close(p[0]); close(p[1]);
wait(NULL); wait(NULL);
```

6 Background Jobs (8) and Process Management

- 8 launches child process **without waiting**.
- Parent returns to shell prompt immediately.
- Must still call `waitpid(-1, &st, WNOHANG)` to reap finished children.

Key functions

Function	Description
<code>fork()</code>	Create new process
<code>execvp()</code>	Replace image with another program
<code>waitpid(pid,&st,WNOHANG)</code>	Reap without blocking
<code>kill(pid,SIGTERM)</code>	Send signal to a process

Zombie: child exits, parent didn't call `wait()` → "defunct" state.

Orphan: parent exits, child adopted by `init/systemd`.

Good practice: install a `SIGCHLD` handler that reaps all children.

7 Threads (Pthreads)

- Threads share address space, heap, globals, and open files.
- Each thread has its own stack and registers.
- Protect shared data with mutexes.

Common functions

Function	Purpose
<code>pthread_create()</code>	Start thread
<code>pthread_join()</code>	Wait for thread
<code>pthread_exit()</code>	End thread explicitly
<code>pthread_detach()</code>	Run thread without joining
<code>pthread_mutex_lock() / unlock()</code>	Protect shared data

Example:

```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
int counter = 0;
void *add(void *arg){
    for(int i=0;i<10000;i++){
        pthread_mutex_lock(&m);
        counter++;
        pthread_mutex_unlock(&m);
    }
    return NULL;
}
```

8 Condition Variables & Synchronization

3 / 7

- Allow threads to sleep until a condition becomes true.
- Always pair with a mutex.
- Use a `while` loop to avoid spurious wakeups.

Wait pattern

```
pthread_mutex_lock(&m);
while(!ready)
    pthread_cond_wait(&cv,&m);
pthread_mutex_unlock(&m);
```

Producer-Consumer pattern

```
pthread_mutex_lock(&m);
while(buffer_full()) pthread_cond_wait(&not_full,&m);
enqueue(item);
pthread_cond_signal(&not_empty);
pthread_mutex_unlock(&m);
```

Deadlock prevention

- Lock in a consistent order.
- Hold locks only briefly.
- Avoid circular dependencies.

9 File Permissions, Ownership & Special Bits

Bit	Applies To	Meaning
r	File/Dir	Read bytes / list directory
w	File/Dir	Write bytes / create-delete entries
x	File/Dir	Execute / traverse directory

Sticky Bit (+t): only file owner or root can delete inside directory (e.g., `/tmp`).
Example: `drwxrwxrwt` → shared temp directory.

10 setuid & setgid (Privilege Bits)

- Purpose:** let users run programs with the owner or group's privileges.
- setuid:** process runs with **EUID = file owner**.
- setgid:** process runs with **EGID = file group**.
- setgid on directories:** new files inherit directory's group.

Notation

4 / 7

Mode	Meaning
-rwsr-xr-x	setuid active (user s)
-rwxr-sr-x	setgid active (group s)
drwxrwsr-x	directory with setgid (shared group)

Examples

- rwsr-xr-x root root /usr/bin/passwd → runs as root.
- drwxrwsr-x root staff /project/shared → files group = staff.

Security

- Works only on binaries.
- Avoid setuid on user-owned scripts.
- Common safe uses: `passwd`, `sudo`, `ping`.

10 Filesystem Inodes & Block Pointers

- Inode** = metadata + pointers to data blocks.
- Pointers are used to access file data stored on disk.

Pointer Type	Count	Capacity (4 KB blocks)
Direct	12	$12 \times 4 \text{ KB} = 48 \text{ KB}$
Single Indirect	1	$1024 \times 4 \text{ KB} = 4 \text{ MB}$
Double Indirect	1	$1024^2 \times 4 \text{ KB} \approx 4 \text{ GB}$
Triple Indirect	1	$1024^3 \times 4 \text{ KB} \approx 4 \text{ TB}$

How to calculate inode capacity

- Direct:** $\# \text{direct} \times \text{block_size}$
→ $12 \times 4 \text{ KB} = 48 \text{ KB}$
- Single indirect:** $\text{block_size} / \text{ptr_size} \times \text{block_size}$
→ $(4096 / 4) \times 4096 = 4 \text{ MB}$
- Double indirect:** $(\text{block_size} / \text{ptr_size})^2 \times \text{block_size}$
→ $(1024^2) \times 4 \text{ KB} = 4 \text{ GB}$
- Triple indirect:** $(\text{block_size} / \text{ptr_size})^3 \times \text{block_size}$
→ $(1024^3) \times 4 \text{ KB} = 4 \text{ TB}$

Path lookup

- Parse path (`/dir1/dir2/file`) → traverse directories.
- Find target inode → read block pointers.
- Access data through direct/indirect blocks.

Links

5 / 7

7 / 7

- Hard link:** another directory entry pointing to same inode.
- Soft link (symlink):** separate inode storing path text.

11 Signals Table

#	Name	Default	Typical Cause / Meaning
1	SIGHUP	T	Terminal hangup / reload config
2	SIGINT	T	Ctrl-C interrupt
3	SIGQUIT	C	Ctrl-\ (quit, core dump)
4	SIGILL	C	Illegal instruction
5	SIGTRAP	C	Debug breakpoint
6	SIGABRT	C	Abort (abort0)
7	SIGBUS	C	Bus error
8	SIGFPE	C	Divide by zero / math fault
9	SIGKILL	T	Uncatchable kill
10	SIGUSR1	T	User-defined
11	SIGSEGV	C	Segmentation fault
12	SIGUSR2	T	User-defined
13	SIGPIPE	T	Write to pipe w/ no reader
14	SIGALRM	T	Timer expired
15	SIGTERM	T	Graceful termination
17	SIGCHLD	I	Child terminated
18	SIGCONT	Cnt	Continue if stopped
19	SIGSTOP	S	Uncatchable stop
20	SIGTSTP	S	Ctrl-Z
28	SIGWINCH	I	Window resize
29	SIGIO	T	I/O possible
30	SIGPWR	T	Power failure
31	SIGSYS	C	Bad syscall

✓ Exam-Day Checklist

- Wire **multi-stage pipelines** with `dup2` and close order.
- Know which process does > `outfile` (**last stage**).

6 / 7

- Explain **FD table vs file table vs inode**.
- Implement a **SIGCHLD reaper** and explain zombies.
- Recall directory **r/w/x** semantics + **sticky bit**.
- Estimate **direct** ($\approx 48\text{KB}$) and **indirect** ($\approx 4\text{MB}$) capacity @ 4KB blocks.
- Condition variable loop: `while (!cond) pthread_cond_wait(...);`